

Death of the Author, Birth of the Scriptor: AI in Academia, Incentives, and Post-Paper Research Outputs^{*}

May 2026[‡]

Abstract

Cesar Hidalgo warns that an AI tsunami is approaching science. The warning matters because AI lowers the execution costs of traditional research tasks: search, coding, modelling, drafting, review, visualisation, and revision. In weak-link terms, AI moves the constraint on research productivity towards specifying, auditing, validating, and warranting the workflow through which claims are produced. This challenges the academic paper as the authoritative way of knowing research. The paper should become the narrative layer of a wider evidence system: a modular constellation of claims, sources, code chains, prompts, checks, exclusions, revisions, uncertainty, and public outputs. Quantity-output governance, patched with ad hoc data, code, and AI-use statements, is poorly suited to rigorous research under generative conditions. Where a prompt uses AI to produce papers, a scriptor recognises that the workflow is the output. **scriptorship-ledger** shows this shift through a claim ledger, transparency log, source-bias report, warrant graph, validation scripts, and workflow profile, paired with a limited v1 Codex/LaTeX pipeline that researches, validates, tracks accepted edits, and makes the writing workflow inspectable.

³Draft dated 13 May 2026.

¹A public workflow repository archives the LaTeX project, claim ledger, transparency log, validation scripts, Overleaf/Git workflow, and prototype reviewer tooling. Codex helped formalise an existing paper-pipeline approach; the initiating prompt is retained as part of the workflow record.

²The title invokes [Barthes \(1977\)](#) to mark a move away from treating the named author as the sovereign origin that fixes meaning. In AI-mediated research, authorial status remains a responsibility category, while claim-level trust has to move through the workflow. For recent discussions of Barthes, authorship, and AI-generated work, see [Masters \(2025\)](#) and [Puscasu \(2024\)](#).

1 Introduction

Cesar Hidalgo’s essay on an AI tsunami facing science is useful because it names a production problem, not just a writing problem (Hidalgo, 2026). AI lowers the execution cost of search, coding, drafting, review, and revision. The immediate risks are familiar: fabricated references, generic prose, weak disclosure, and uncertain responsibility (Committee on Publication Ethics, 2023; Masters, 2025; Puscasu, 2024). But those failures do not exhaust the issue. As agentic systems improve, the harder case is plausible research whose path from source to claim is opaque (Lu et al., 2026; Brodeur et al., 2026). RoRI’s point that no paper is perfect helps locate the unit of repair: claims, not whole papers (Research on Research Institute, 2026). The question is therefore not whether researchers can ask better questions while keeping the paper unchanged. It is whether the paper remains a suitable object of trust under accelerated, modular production. This article argues for workflow evidence, claim ledgers, and scriptorship.

2 Output-Based Trust

Academic assessment runs on proxies. Papers, journal placement, author status, institution, citation counts, h-indexes, publication volume, and REF-style outputs stand in for quality. The main problem with current papers is not the paper form itself, but its institutionalisation as the privileged unit of account. Paper publishing became central to hiring and promotion because it gave institutions visible signals of productive research when direct workflow inspection was costly.

The proxies were already strained. DORA, CoARA, the Leiden Manifesto, and the Hong Kong Principles all argue against inappropriate reliance on journal status and publication metrics (DORA, 2013; Coalition for Advancing Research Assessment, 2022; Hicks et al., 2015; Moher et al., 2020). REF 2029 still organises output eligibility and open access around publication categories, even while recognising other outputs such as datasets and code outside the core open-access requirement (Research Excellence Framework, 2025). Publication bias, false-positive risk, citation inflation, gatekeeping, high journal costs, and reviewer overload all predate current AI systems (Ioannidis, 2005). The result is a bloated quantity-output system: built to count, rank, and process outputs, not to inspect the workflow by which claims become trustworthy.

The paper should not be treated as the authoritative way of knowing research when evidentiary work is distributed across sources, scripts, prompts, reviews, revisions, and checks. Nor is it enough to leave the paper in place and attach ad hoc data availability, code, or AI-use statements. The shift should be more systematic: from outcome-centred governance to a workflow-based model in which the paper becomes the narrative layer of a wider evidence system. The research of the future will be scripted, not merely scribed.

Current academic incentives work in the opposite direction. Gartenberg and colleagues argue that AI enters an institutional setting already organised by publish-or-perish pressures, submission growth, reviewer overload, and quantity-oriented assessment (Gartenberg et al., 2026). Their more-versus-better argument matters because it treats AI-generated research as an institutional incentive problem, not simply a tool-use problem. Big data and AI create real potential for better research, but that potential is weakened if institutions reward faster production of polished paper outcomes rather than more trustworthy workflows. Governance remains attached to the paper, the author, and the publication event, while AI-driven research increasingly happens through modular, iterative, tool-mediated workflows.

COPE guidance on AI authorship, data availability statements, code-sharing policies, and AI disclosure statements remain necessary but insufficient for an era in which polished research outputs

become easier to produce. Generative tools expose this weakness by making more outputs easier to generate, revise, and circulate.

Output governance asks who authored the paper, whether AI was disclosed, whether there is a code link, and whether the article is suitable for publication. Workflow governance asks a different set of questions: what produced this claim, what sources entered, what was excluded, what scripts transformed the evidence, where humans audited, which specifications were tried, which checks failed, and whether the claim can be inspected, reproduced, challenged, and repaired. The aim is a workflow-based assessment standard, not another cumulative scoreboard: a way to privilege trust production over paper production.

Existing governance is incomplete. A data availability statement matters only when the data connect to the claims. A code link matters when scripts reproduce the relevant tables, figures, and transformations. An AI disclosure matters when it distinguishes drafting, search, coding, review, and judgement. A journal decision matters when venue status does not replace claim-level warrant.

3 Falling Execution Costs

Hidalgo’s essay uses William Shockley’s model of research productivity to frame scholarly work as a set of complementary skills (Hidalgo, 2026). In weak-link or O-ring logic, output is constrained by the least adequate component. A researcher may be held back by coding, plotting, literature search, writing, mathematics, or review. If AI systems reduce the cost of some tasks, the binding constraint moves.

The Nature AI Scientist paper illustrates the movement. Its system performs ideation, literature search, code generation, experimentation, plotting, manuscript drafting, and automated review (Lu et al., 2026). The authors report underdeveloped ideas, implementation errors, hallucinated citations, and inconsistent performance. The mechanism still matters: some components of research execution are becoming cheaper, faster, and more iterable.

If execution becomes cheaper, specification search becomes cheaper. Researchers can try more models, prompts, drafts, robustness checks, source combinations, and reviewer loops. This expands the possibility set of research. It can advance understanding when the workflow is disciplined, logged, and audited. It can also expand the garden of forking paths (Gelman and Loken, 2013). The many-analysts literature shows that analytic choices can matter even when teams start from the same data and question (Silberzahn et al., 2018).

The weak link therefore shifts towards scripting, verification, judgement, source selection, and warranted trust. Workflow integrity becomes central to filtering the noise of AI-accelerated output. AI may simply make the treadmill faster. Workflow evidence makes acceleration inspectable.

Tool sophistication extends beyond reviewer bots. Agents can already operate inside command-line environments, sandboxes, browsers, repositories, documentation search tools, and API-connected workflows. They can pull an Overleaf repository, inspect diffs, run tests, build figures, query documentation, maintain logs, and check references against a structured source ledger. The institutional question is whether those capacities multiply outputs or expose the path from source to claim.

4 Birth of the Scriptor

Authorship is no longer a sufficient description of research production. The author remains accountable, but responsibility shifts into workflow design, validation, and auditability. A human researcher remains accountable for what is asked, accepted, rejected, checked, revised, and published.

The visible intellectual labour now also includes designing the workflow through which sources, code, prompts, models, reviews, revisions, and public outputs interact.

AI authorship is disingenuous when it merely supplies outputs to an output-focused incentive system. AI scriptorship can be productive when it makes research workflows inspectable, auditable, and claim-level. A prompter asks for a result: write a paragraph, code a model, summarise a paper, generate a table. Where a prompter uses AI to produce papers, a scriptor recognises that the workflow is the output. The scriptor designs the conditions under which results are produced, checked, rejected, revised, and exposed.

This distinction shifts attention from unassisted human writing to trustworthy human understanding. A research output should further understanding, and the workflow behind it should be transparent enough to deserve trust. In a world of rapidly falling execution costs, the weak link is scripting.

Transparency must shift accordingly. It is no longer enough to disclose whether a named author wrote the final text or whether a coder wrote a script by hand. Readers need to know what was asked, generated, accepted, rejected, checked, changed after review, and warranted by evidence.

The scriptor gives a stricter account of authorship. The researcher is responsible for sentences and for the workflow that made those sentences plausible: refusing attractive outputs, preserving uncertainty, maintaining source provenance, and excluding generated results that should stay outside the scholarly record. Scriptorship names that responsibility.

5 Modularisation of Research and Knowledge

AI-mediated research production is modular. Sources, prompts, code, data transformations, figures, tables, drafts, reviews, revisions, appendices, repositories, and public essays can be produced and revised semi-independently. The paper remains the coherent narrative. It is no longer the whole object. Research output becomes a constellation: article, ledger, code chain, source log, warrant graph, review history, appendix, and public account.

Current responses to AI-generated papers show the limit of paper-centred adaptation. The Journal for AI Generated Papers acknowledges AI-generated submissions, screening, endorsement, AI review, and future peer review ([Journal for AI Generated Papers, 2026](#)). Yet the paper remains the unit of account. A constellation approach asks whether that unit of account is adequate for a production system built from sources, prompts, code, models, reviews, revisions, and validation traces. The relevant question is what combination of paper, source log, code chain, prompt record, claim ledger, review trace, and warrant graph best advances understanding.

Agentic coding and CLI research agents make this modularity operational. A command-line agent can pull a repository, edit a script, run a test, inspect an error, search documentation, update a source log, and propose a prose revision as separable steps. Trust still requires audit, but the chain of knowledge production becomes easier to expose: which instruction produced which file, which check failed, which diff was accepted, and which claim changed as a result ([Hidalgo, 2026](#); [Lu et al., 2026](#); [Vincent and contributors, 2026](#)). The relevant skill is designing human-audited workflows.

Knowledge claims are modular too. A paper can contain a strongly warranted replication, a moderately warranted interpretation, a weak speculative extension, and a conceptual claim that is not directly testable. Treating the whole paper as simply accepted or rejected, good or bad, human or AI-assisted, is too blunt.

The 2026 Nature reproducibility and robustness study should not be overread. It does not show that most economics and political science papers are not robust. Among 110 articles from leading journals with mandatory data and code-sharing policies, more than 85% of published claims were

computationally reproducible and 72% of statistically significant estimates remained significant and in the same direction under robustness checks (Brodeur et al., 2026). The narrower lesson is that even relatively good transparency conditions still require workflow-level assessment.

PDF appendices help, but they freeze modular evidence into a static form. Online appendices are better, but often remain hard to interrogate. Modular research production requires modular knowledge assessment. This motivates claim-level warranting.

The modularity is also social. CRediT-style contribution taxonomies already recognise that research is produced through differentiated roles rather than a single undivided authorial act (NISO, 2022). Research objects and FAIR principles similarly treat digital scholarly materials as things that should be findable, accessible, interoperable, and reusable (Bechhofer et al., 2010; Wilkinson et al., 2016). Scriptorship applies that logic to AI-mediated workflows and to the claims those workflows produce.

6 The SCRIPT Protocol

SCRIPT is a minimal protocol for workflow evidence. It defines what a research output includes when AI systems, APIs, scripts, repositories, browser-capable agents, RAG indexes, and long-running background tasks enter research production.

S is Sources: provenance, search transparency, sources searched and screened, screening depth, sources cited and excluded, source type, discipline, geography, language, date bias, known blind spots, and selection bias in the literature itself.

C is Code chain: how raw materials become claims, including Git history, LaTeX build, Overleaf/Git sync, reproducible scripts, git-diff review, API integrations, and agentic execution logs.

R is Researcher/AI roles: what humans did, what AI did, what was rejected, what was checked, where judgement entered, and how author, prompter, scriptor, reviewer, auditor, and maintainer roles differ. Human responsibility remains central.

I is Interactive inspection: appendices that allow readers and future agents to ask which claim a source supports, what was excluded, what scripts generated a table, what uncertainty attaches to a claim, what a reviewer flagged, and what changed after review.

P is P-hacking visibility: model choices, tested alternatives, exclusions, transformations, outcomes, control sets, robustness checks, and stopping rules. P-hacking should be findable in the scripts.

T is Tests: validation, robustness, missingness, sensitivity, reproducibility, overclaiming checks, causal-language checks, source-bias checks, and warrant-graph consistency checks.

SCRIPT has to be light enough to use. A protocol that requires every researcher to build a bespoke infrastructure stack will fail. The minimum version is a small set of structured files and scripts: sources, claims, roles, logs, checks, and uncertainty. More demanding projects can add executable environments, containers, notebooks, browser inspection, or RAG interfaces. The first requirement is simple: the workflow must be explicit enough to inspect.

7 Claim-Level Warranting

The claim is a better unit of evaluation than the whole paper. A paper may contain a strongly warranted empirical claim, a moderately warranted descriptive claim, a weak speculative claim, and a conceptual claim that is not directly testable. Readers should see that structure without inferring it from prose alone.

Warrant should be non-additive and weak-link based. A claim cannot be strongly warranted if one essential evidentiary link is opaque. A beautiful figure cannot compensate for an unverifiable

data transformation. A long reference list cannot compensate for a citation that does not support the sentence it is attached to. A transparent repository cannot compensate for causal language attached to a descriptive design.

This differs from metric accumulation. Credit should be earned by making claims more trustworthy. The h-index counts the afterlife of papers. A warrant graph inspects the life of claims.

The warrant graph is a map of dependencies. A reader can ask why a conceptual claim is moderate rather than strong, why a speculative claim remains weak, or why an empirical claim depends on a particular source or script. Disagreement is expected. The value lies in making disagreement local. Instead of rejecting a whole paper because one claim is overdrawn, reviewers can repair the claim, narrow it, or attach uncertainty.

8 Verification

Verification should let readers, reviewers, and future agents inspect enough of the system without relying on authorial status alone. They should be able to examine the claim ledger, source log, warrant graph, overclaiming checks, source-bias report, transparency log, Overleaf/Git diffs, AI/human role statement, and validation scripts.

Open science frameworks already point this way. TOP Guidelines, FAIR principles, research objects, software citation principles, CRediT, and Registered Reports all make parts of the research process more inspectable, reusable, and creditable (Nosek et al., 2015; Wilkinson et al., 2016; Bechhofer et al., 2010; Smith et al., 2016; NISO, 2022; Nosek and Lakens, 2014). Scriptorship extends this logic to AI-mediated workflow governance.

The tools reach beyond reviewer bots. They can be integrated into Overleaf/Git writing, literature review, source logging, claim extraction, style and voice preservation, workflow validation, and revision loops. CLI agents, sandbox agents, browser-capable agents, MCP-style documentation search, API integrations, RAG indexing, and multi-step tool use can all participate in this workflow.

Verification is partial by design. A repository makes the path to a claim easier to inspect. A reader may verify a source-to-claim link, rerun a validation script, inspect a rejected output, check whether a limitation was acknowledged, or challenge a warrant level. The quality standard must shift from polished outputs to inspectable workflows.

This changes what appendices are for. The appendix becomes a verification interface. A static table may still be needed for archival stability, but readers should also be able to query the evidence system. Future appendices should support claim lookup, source lookup, exclusion lookup, specification lookup, review-history lookup, and uncertainty lookup. The PDF remains the narrative; the repository becomes the inspection layer.

9 Incentives

Paper publishing became tied to promotion because, for a long time, papers, venues, and citations were rough but workable signals of productive research. AI changes that production function. As tools become more capable at drafting, coding, reviewing, revising, and packaging, polished output production becomes radically faster and more widely accessible. The signal carried by paper volume therefore weakens. Organization Science’s 2026 editorial reports a 42% rise in submission volume since the release of ChatGPT and argues that current tools, amplified by publish-or-perish incentives, are pushing the system towards more rather than better research (Gartenberg et al., 2026). Bloating and lack of clarity are failures of research as an activity organised around understanding.

The rise of AI-generated papers (AIGPs) therefore exposes a lag in publication-based incentive structures. If credit continues to flow through paper counts, venue prestige, citation accumulation, and impact narratives attached to outputs, AI will be used to produce more countable artefacts. That incentive structure also rewards ambiguity. When disclosure threatens the status of a finished paper, researchers have reason to hide, soften, or minimise AI use rather than expose it productively. University discussions of impact will then be limited to unproductive AI masking, rather than a more expansive use of AI through scriptorship. If credit flows through workflow integrity, validation, and repair, AI can instead be directed towards making claims more trustworthy. The margin of skill moves from producing a polished artefact to designing and auditing the workflow that makes the artefact reliable.

This is also an attention problem. In a saturated publication system, research is often encountered through titles, abstracts, search results, citation networks, dashboards, and summaries before it is read in full. AI-generated volume makes that noise worse. Workflow evidence helps readers shift from output polish to inspectable warrant.

Review should also become claim-level. A good review identifies the affected claim, weak link, evidence, severity, and repair path. This avoids two failures: destructive whole-paper criticism, where every imperfection becomes fatal, and permissive acceptance, where fluent prose passes despite opaque claims. The RoRI framing that no paper is perfect helps because review should improve warrant without pretending that perfection is available ([Research on Research Institute, 2026](#)).

Review should become two-way. Authors gain because the work becomes stronger. Reviewers gain because their epistemic contribution becomes visible. A sampled warrant auditor may provide a third role, checking whether claims, sources, and validation evidence align without creating infinite regress.

A reviewer who improves the warrant of a claim has made a scholarly contribution, even if the final paper is not theirs. Current systems hide that contribution inside anonymous labour. A workflow-governed system could credit review that identifies a weak source, repairs an overclaim, improves a test, or clarifies a limitation. The aim is repairable criticism.

10 Authorship, Voice, and Workflow Transparency

The aim is trustworthy human understanding. AI tools can help draft and revise, but responsibility lies in specification, selection, rejection, verification, and judgement.

Voice matters because homogenised prose can blur distinctions, hide uncertainty, and weaken conceptual force. Accepted human edits should be treated as evidence of judgement.

Git-diff voice learning provides one route. The researcher edits text directly in Overleaf. The repository pulls those edits through Git. The system analyses the diff, extracts before/after examples, and proposes voice or style decision rules for later prose passes. These rules may preserve British English, reduce defensive caveats, avoid clunky self-reference, or protect strong conceptual lines. They never automatically overwrite prose.

This makes voice part of the workflow rather than a mystical residue. The system learns from accepted human edits while preserving the author's ability to reject the rule. The goal is to preserve authorial judgement while making the workflow inspectable.

Voice learning also guards against thin transparency. A paper can disclose AI use and still sound as if its sentences were flattened by a generic assistant. A workflow can show how the human author shaped the argument through accepted edits. The diff records judgement: refusal, sharpening, narrowing, and revoicing.

11 Implementation: scriptorship-pipeline and scriptorship-ledger

scriptorship-pipeline is the reusable tooling layer: Overleaf/Git integration, git-diff review, git-diff voice learning, claim extraction, source logging, source-bias reports, overclaiming checks, warrant graphs, validation scripts, RAG appendix specification, and agent rules. In its current v1 form, it is a Codex integration for LaTeX/Git work: it operates directly on the article files, researches and logs sources, validates structured evidence, checks references and claims, tracks the author’s accepted changes through Git diffs, and turns those changes into proposed future prose rules.

scriptorship-ledger contains the LaTeX article, Substack draft, claim ledger, source ledger, transparency log, source-bias report, warrant graph, SCRIPT framework, uncertainty register, workflow profile, time-savings estimate, AI-use statement, human-audit statement, prompt log, Overleaf notes, and appendices.

The implementation borrows from an existing paper-pipeline architecture: a hub rather than scratch sprawl; Overleaf checkouts as Git repositories; release gates; citation and source robustness audits; reviewer-failure logic; run summaries; verification JSON; and style learning from diffs. Obra’s Superpowers provide a reference point for modular agent skills, explicit rules, TDD, review, and verification habits (Vincent and contributors, 2026).

The current **scriptor-pipeline-v1** implementation works with Superpowers in Codex as a small demonstration of that shift: skills and agent rules organise the work, while the pipeline records claims, sources, checks, and warrant evidence. Its limitations are part of the point. Even a lightweight prototype shows that writing, researching, reference checking, revision, validation, and public accounting can be treated as one inspectable workflow rather than as a private production process followed by a finished PDF.

The prototype uses Python, JSON-valid YAML, Markdown, LaTeX, and Mermaid rather than a paid API or hidden service. A laboratory project, an ethnography, a computational social-science paper, and a theory article will need different checks. The common point is inspectable provenance and explicit warrant. That implies a ledger standard alongside the paper: a durable record of sources, prompts, scripts, accepted edits, rejected outputs, validations, uncertainty, and claim-level warrant.

12 Conclusion

A bloated quantity-output system is poorly suited to rigorous research in a world where polished outputs become cheap. A fluent paper can be weakly warranted. A clumsy paper can contain a strong claim. An AI-assisted workflow can be opaque. A human-written workflow can be opaque too.

AI authorship in output-focused systems feeds ambiguity back into publication-based incentives. AI scriptorship in workflow-focused systems can be auditable. The researcher’s role is auditing the workflow through which claims are produced.

The article and analysis remain important. They organise thought, make judgement legible, and give readers a narrative path through evidence. But they should be balanced within a workflow that recognises how AI-generated material can enter the production of understanding. Sources, code chains, roles, appendices, specification paths, tests, review interventions, and warrant graphs should be part of what is assessed.

Workflow integrity and diligence become central in filtering the growing noise associated with AI-accelerated research outputs. The quality standard must shift from polished outputs to inspectable

workflows because the existing machinery is not built to handle these generative trends. Without that shift, AI becomes a problem to hide at the publication boundary rather than a tool whose contribution to understanding can be organised, tested, and credited. Credit should be earned by making claims more trustworthy.

The author remains responsible, but authorial status moves away from the centre of trust. What matters is the governed production of understanding. The scriptor turns skill sets into human-audited workflows others can inspect.

References

- Roland Barthes. The death of the author. In Stephen Heath, editor, *Image, Music, Text*, pages 142–148. Fontana Press, London, 1977. Original essay published in French in 1967.
- Sean Bechhofer, David De Roure, Matthew Gamble, Carole Goble, and Iain Buchan. Research objects: Towards exchange and reuse of digital knowledge, 2010.
- Abel Brodeur, Derek Mikola, Nikolai Cook, Lenka Fiala, Thomas Brailey, Ryan Briggs, et al. Reproducibility and robustness of economics and political science research. *Nature*, 652: 151–156, 2026. doi: 10.1038/s41586-026-10251-x. URL <https://www.nature.com/articles/s41586-026-10251-x>.
- Coalition for Advancing Research Assessment. Agreement on reforming research assessment, 2022. URL <https://www.coara.org/agreement/the-agreement-full-text/>.
- Committee on Publication Ethics. Authorship and ai tools, 2023. URL <https://publicationethics.org/cope-position-statements/ai-author>.
- DORA. San francisco declaration on research assessment, 2013. URL <https://sfdora.org/about-dora/>.
- Claudine Gartenberg, Sharique Hasan, Alex Murray, and Lamar Pierce. More versus better: Artificial intelligence, incentives, and the emerging crisis in peer review. *Organization Science*, 2026. doi: 10.1287/orsc.2026.ed.v37.n3. URL <https://pubsonline.informs.org/doi/10.1287/orsc.2026.ed.v37.n3>.
- Andrew Gelman and Eric Loken. The garden of forking paths, 2013. URL http://www.stat.columbia.edu/~gelman/research/unpublished/p_hacking.pdf.
- Diana Hicks, Paul Wouters, Ludo Waltman, Sarah de Rijcke, and Ismael Rafols. Bibliometrics: The leiden manifesto for research metrics. *Nature*, 520:429–431, 2015. doi: 10.1038/520429a.
- Cesar Hidalgo. An ai tsunami is about to hit science, March 2026. URL <https://cesarhidalgo.com/blog/2026/3/6/an-ai-tsunami-is-about-to-hit-science>. Accessed 12 May 2026.
- John P. A. Ioannidis. Why most published research findings are false. *PLOS Medicine*, 2(8):e124, 2005. doi: 10.1371/journal.pmed.0020124.
- Journal for AI Generated Papers. Journal for ai generated papers, 2026. URL <https://jaigp.org/>. Accessed 12 May 2026.
- Chris Lu, Cong Lu, Robert Tjarko Lange, Yutaro Yamada, Shengran Hu, Jakob Foerster, David Ha, and Jeff Clune. Towards end-to-end automation of ai research. *Nature*, 651:914–919, 2026. doi: 10.1038/s41586-026-10265-5. URL <https://www.nature.com/articles/s41586-026-10265-5>.
- Ken Masters. Artificial intelligence and the death of the academic author. *Medical Teacher*, 47(10): 1563–1565, 2025. doi: 10.1080/0142159X.2025.2523463.
- David Moher, Lex Bouter, Sabine Kleinert, Paul Glasziou, Mai Har Sham, Virginia Barbour, et al. The hong kong principles for assessing researchers: Fostering research integrity. *PLOS Biology*, 18(7):e3000737, 2020. doi: 10.1371/journal.pbio.3000737.

- NISO. Contributor role taxonomy, 2022. URL <https://credit.niso.org/>.
- Brian A. Nosek and Daniel Lakens. Registered reports: A method to increase the credibility of published results. *Social Psychology*, 45(3):137–141, 2014. doi: 10.1027/1864-9335/a000192.
- Brian A. Nosek, George Alter, George C. Banks, Denny Borsboom, Sara D. Bowman, Steven J. Breckler, et al. Promoting an open research culture. *Science*, 348(6242):1422–1425, 2015. doi: 10.1126/science.aab2374.
- Iulia-Dana Puscasu. Creating with ai: On recent debates about authorship revisiting the influence of barthes and foucault. *Studia Universitatis Babes-Bolyai Philosophia*, 69(Special Issue):49–60, 2024. doi: 10.24193/subbphil.2024.sp.iss.03.
- Research Excellence Framework. Ref 2029 open access policy, 2025. URL <https://2029.ref.ac.uk/guidance/ref-2029-open-access-policy/>.
- Research on Research Institute. No paper is perfect, 2026. URL <https://researchonresearchinstitute.substack.com/p/no-paper-is-perfect>. Accessed 12 May 2026.
- Raphael Silberzahn, Eric Luis Uhlmann, Daniel P. Martin, Pasquale Anselmi, Frederik Aust, Eli Awtrey, et al. Many analysts, one data set: Making transparent how variations in analytic choices affect results. *Advances in Methods and Practices in Psychological Science*, 1(3):337–356, 2018. doi: 10.1177/2515245917747646.
- Arfon M. Smith, Daniel S. Katz, Kyle E. Niemeyer, and FORCE11 Software Citation Working Group. Software citation principles. *PeerJ Computer Science*, 2:e86, 2016. doi: 10.7717/peerj-cs.86.
- Jesse Vincent and contributors. Superpowers, 2026. URL <https://github.com/obra/superpowers>. Accessed 12 May 2026.
- Mark D. Wilkinson, Michel Dumontier, IJsbrand Jan Aalbersberg, Gabrielle Appleton, Myles Axton, Arie Baak, et al. The fair guiding principles for scientific data management and stewardship. *Scientific Data*, 3:160018, 2016. doi: 10.1038/sdata.2016.18.