

Qualification Structure and Local Employment Rates, 2016–2025

Nicholas P. Sweeney

May 18, 2026*

Abstract

This paper studies whether pre-2020 NVQ4 qualification share predicts employment rate in a NOMIS Annual Population Survey panel. The analysis retains 3,800 unit-year observations for 380 units from 2016 to 2025. The sample spans 2016–2025, giving 5 post-2020 years for a dynamic adjustment reading. The dynamic estimate at event time 5 is -1.180 (SE 0.608), and the controlled post-2020 fixed-effects coefficient is 0.463 (SE 0.248). The specifications include unit and year fixed effects, unit-clustered standard errors, and baseline controls interacted with the post period: pre-2020 employment rate, pre-2020 economic activity rate, pre-2020 unemployment rate. The contribution is a human-capital adjustment estimate that links the coefficient path to a named local mechanism and a source-specific control set.

1 Introduction

Skill composition matters for local employment because it affects job matching, occupational mobility, and the speed with which workers can shift across firms and sectors (Busso et al., 2013; Angrist and Pischke, 2009). This paper asks: Do local authorities with higher qualification shares have different employment-rate paths in the APS panel? It studies 3,800 UK local authority-year observations for 380 units over 2016–2025, using NOMIS Annual Population Survey and a local-authority fixed-effects APS panel (for National Statistics, 2026; Kitchin, 2014).

The issue is comparison quality (Greenstone et al., 2010; Rosenbaum, 2002). Pre-2020 nvq4 qualification share is economically meaningful, but it also sits beside baseline differences in scale, density, skills, sector mix, and labour-market status. The paper therefore estimates the qualification-gradient employment path with unit and year fixed effects and post-period interactions for pre-2020 employment rate, pre-2020 economic activity rate, pre-2020 unemployment rate (for National Statistics, 2026; Peng, 2011; Wilkinson et al., 2016).

The benchmark is the human-capital and local-labour-market literature, where skill composition matters because it shapes occupational mobility, employer matching, and the local response to demand shifts. The design follows that standard by tying a distinct source, question, outcome, exposure, control set, and robustness margin to one empirical mechanism (Kitchin, 2014; Borgman, 2015).

The main result is interpretable in economic terms (Austin et al., 2018; Arkhangelsky et al., 2021). The event-time 5 coefficient is -1.180 (SE 0.608); the compact post-period coefficient is 0.463 (SE 0.248). A positive estimate means high-qualification local authorities sustain higher employment rates after the event; a negative estimate means qualification structure is not translating into stronger employment-rate adjustment in this panel (Bureau, 2025; Wilkinson et al., 2016).

*Manuscript date: May 18, 2026.

The example is concrete (Chodorow-Reich, 2019; Imbens and Rubin, 2015). The exposure is fixed before the post-period, the outcome is observed repeatedly for the same units, and the controls are chosen because the literature treats them as baseline channels rather than incidental labels. That structure lets the paper separate a capacity-gradient estimate from a loose narrative about places doing better or worse (Callaway and Sant’Anna, 2021; Sun and Abraham, 2021).

The comparison is deliberately narrow. A distinct source, outcome, exposure, control set, and robustness margin keep the claim tied to one economic object rather than to title variation alone (Goodman-Bacon, 2021; Arkhangelsky et al., 2021).

The issue, then, is not whether the paper can sound plausible (Bartik, 1991; Sun and Abraham, 2021). The issue is whether the empirical object deserves interpretation. The example here is pre-2020 NVQ4 qualification share: it has a clear pre-period definition, a local economic mechanism, and controls that speak to the outcome process. The analysis reads the coefficient only after those elements are in place (Abadie et al., 2010; Athey and Imbens, 2016).

The contribution is a disciplined estimate rather than a broad story (Kline and Moretti, 2014; Gelman et al., 2020). It gives the reader a panel, a figure, a regression table, a robustness table, and an audit file, all tied to the same question. That package is what lets a reviewer identify the next missing variable or stronger design without first repairing the manuscript structure (Chernozhukov et al., 2018; Wager and Athey, 2018).

Section 2 discusses the local adjustment literature; Section 3 defines the panel, exposure, outcomes, and controls; Section 4 sets out the fixed-effects design; Section 5 reports the event path and regression estimates; and Section 6 relates the evidence back to the mechanism, identification limits, and next evidence.

2 Literature

2.1 Local Multipliers

The human-capital literature motivates the question, while the local-labour-market literature warns that employment, activity, and unemployment baselines must travel with the skill exposure (Moretti, 2011; Glaeser, 2008; Autor et al., 2013).

The example in this paper is pre-2020 NVQ4 qualification share. It is a useful pre-period margin because it gives the post-period comparison an economic object: employment adjustment through local human-capital composition. The relevant question is whether that object organizes employment rate once standard baseline controls are allowed to matter (Bartik, 1991; Kline and Moretti, 2014).

Controls are part of the theory rather than table dressing. Population scale changes market size; density changes land, commuting, and matching conditions; education and qualifications change the skill margin; baseline employment, unemployment, and activity rates describe local labour-market slack. Those variables belong in the design when the outcome calls for them (Busso et al., 2013; Greenstone et al., 2010).

The identification literature raises the harder question: were high-exposure places already moving differently? The event-study plot keeps that question visible by showing pre-period coefficients, and the pre-trend placebo reports the same concern in a compact regression row (Neumark and Simpson, 2015; Austin et al., 2018).

2.2 Identification

The contribution is focused: a public-data estimate of employment adjustment through local human-capital composition. The paper adds value when the coefficient path, control set, and robustness margin all point to a coherent interpretation (Moretti, 2011; Glaeser, 2008).

The highest bar in related applied work is set by papers that join mechanism, measurement, and comparison design (Kline and Moretti, 2014; Gelman et al., 2020). Large-plant, place-based policy, local-labour-market, and spatial-equilibrium studies each make the same demand in different ways: the local variable must have an economic role, and the empirical design must show why the comparison is informative (Greenstone et al., 2010; Busso et al., 2013; Kline and Moretti, 2014; Moretti, 2011; Glaeser, 2008; Monte et al., 2018).

That benchmark also shapes the diagnostic standard. Difference-in-differences papers now report timing, treatment heterogeneity concerns, and pre-period checks; observational studies report sensitivity and design assumptions; spatial papers worry about spillovers, clustering, and the modifiable areal unit problem (Angrist and Pischke, 2009; Wooldridge, 2010; Callaway and Sant’Anna, 2021; Sun and Abraham, 2021; Goodman-Bacon, 2021; Rosenbaum, 2002; Conley, 1999; Cameron et al., 2011; Openshaw, 1984).

The control set follows that logic (Greenstone et al., 2010). For U.S. county papers, log population, density, education, and sector shares describe market scale, land intensity, human capital, and composition. For NOMIS papers, pre-period employment, unemployment, economic activity, and qualifications describe labour-market attachment and skill structure. The controls therefore match the outcome process rather than merely filling a regression table (Athey and Wager, 2019; Mullainathan and Spiess, 2017).

The methods frontier is useful as a discipline even when the final model is simpler (Neumark and Simpson, 2015; Manski, 2003). Synthetic controls and synthetic difference-in-differences require donor support; partial identification is valuable when treatment or spillovers are interval-valued; double machine learning and causal forests require richer covariates and a reason to study heterogeneity (Abadie et al., 2010; Arkhangelsky et al., 2021; Manski, 2003; Chernozhukov et al., 2018; Wager and Athey, 2018; Athey and Wager, 2019; Athey and Imbens, 2016).

The paper also follows reproducible-research standards. The public source is named, raw payloads are hashed, derived files are retained, and the quality audit records the benchmark, pass conditions, remaining development need, and next upgrade (Peng, 2011; Wilkinson et al., 2016; Kitchin, 2014; Borgman, 2015).

A useful literature review also disciplines interpretation (Chodorow-Reich, 2019). If the outcome is employment rate, the paper should explain why that outcome is the right margin for employment adjustment through local human-capital composition. If the robustness outcome is unemployment rate, the paper should say whether it is a mechanism check, a neighbouring outcome, or a stress test on measurement. This framing keeps the result tied to economic content (Gentzkow et al., 2019; Grimmer and Stewart, 2013).

The best papers are also clear about what would change their reading (Moretti, 2011; Rosenbaum, 2002). A strong post-period coefficient with a weak pre-trend looks different from a post-period coefficient that continues an old drift. A robustness result that agrees with the main outcome looks different from one that breaks the mechanism. This paper carries those distinctions into the results section rather than leaving them implicit (Gelman and Hill, 2007; Gelman et al., 2020).

3 Data

3.1 Source and Sample

The Annual Population Survey supplies local-authority employment, unemployment, activity, and qualification measures. The exposure is the pre-period NVQ4 share; the outcome is the employment rate; the robustness margin is unemployment. The derived sample contains 3,800 observations, 380 units, and the 2016–2025 panel window ([for National Statistics, 2026](#); [Kitchin, 2014](#)).

The exposure is measured before 2020. Units above the median pre-period value are classified as high exposure, and the continuous-exposure model keeps the underlying gradient in the table. The outcome is employment rate; the robustness outcome is unemployment rate ([Duranton and Venables, 2015](#); [Fujita et al., 1999](#)).

The controls are pre-2020 employment rate, pre-2020 economic activity rate, pre-2020 unemployment rate. They are reported with readable labels in the summary table and written to the derived panel so the empirical object can be inspected outside the manuscript.

Provenance is handled as a data constraint. Raw public payloads are downloaded, hashed, transformed into derived CSVs, and discarded. The retained files include the panel, event-study estimates, regression estimates, robustness checks, download manifest, and critical audit ([Iammarino et al., 2019](#); [Rodriguez-Pose, 2018](#)).

3.2 Exposure and Outcomes

The main measurement risk is specific to the source. APS local-authority estimates are survey based, so small-area noise and changing samples matter for interpretation. That risk guides the interpretation and identifies the most valuable next data join.

The sample construction is intentionally strict ([for National Statistics, 2026](#)). Units enter the analysis only when the pre-period exposure can be computed, which keeps the high-exposure indicator anchored before the post-period outcome. Years enter the model sample through the profile window, so a paper with too short a post-period is blocked before it can make a broad temporal claim ([Kitchin, 2014](#); [Borgman, 2015](#)).

The table labels are written in economic language ([Kitchin, 2014](#); [Greenstone et al., 2010](#)). The generated summary table reports controls such as pre-2020 employment rate, pre-2020 economic activity rate, pre-2020 unemployment rate; the regression table reports high exposure by post-period and continuous exposure rows; the robustness table names unemployment rate. Those labels are chosen for readers, not for internal variable names.

The derived panel also supports direct inspection ([Borgman, 2015](#); [Peng, 2011](#)). A reader can open the CSV, see unit identifiers, years, outcomes, exposure, high-exposure status, event time, and controls, then compare those fields to the event-study and regression outputs. That path is important because the prose should interpret a visible empirical object rather than ask for trust ([Anselin, 1988](#); [LeSage and Pace, 2009](#)).

Missingness and aggregation are treated as part of measurement ([for National Statistics, 2026](#)). Public sources differ in frequency, geography, and small-area reliability; the generator therefore uses annual unit panels and profile-specific controls rather than forcing every source into the same data shape. The result is comparable across papers without pretending that CBP counties, Census migration counties, NOMIS claimant authorities, and APS survey estimates are the same object ([Openshaw, 1984](#); [Fotheringham et al., 2002](#)).

The outcome scale is also stated directly ([Kitchin, 2014](#); [Chodorow-Reich, 2019](#)). Rates, growth rates, survey percentages, and log-count transformations have different economic meanings. The manuscript therefore reports the outcome label in the abstract, summary table, regression table, and

results discussion, so the reader sees whether a coefficient is a rate-point movement, a growth-rate contrast, or a log-count difference (Pearl, 2009; Keele, 2015).

The data audit records the same principle in file form (Borgman, 2015; Peng, 2011). It states the benchmark paper standard, the pass conditions for panel support, the controls, the timing evidence, the static coefficient, the robustness evidence, and the remaining development need. That file matters because it turns quality into something inspectable: a future run can improve the missing mechanism data without rediscovering what the present panel already establishes (King et al., 1995; Shadish et al., 2002).

The audit also keeps the data and prose synchronized. If the panel window, controls, or robustness outcome change, the title, abstract, tables, and discussion change with them. That prevents a paper from carrying an old frame after the data have moved (Peng, 2011; Wilkinson et al., 2016).

Table 1: Downloaded NOMIS Annual Population Survey panel

Quantity	Value
Unit-year rows retained	3800
Units	380
Years	2016–2025
Exposure threshold	37.100
Controls	pre-2020 employment rate, pre-2020 economic activity rate, pre-2020 unemployment rate
Raw payloads retained	No

Notes: Controls are baseline or pre-period measures used as post-period interactions in the fixed-effects specifications.

4 Empirical Strategy

4.1 Event Study

The event-study specification estimates dynamic high-exposure differences in employment rate around 2020:

$$Y_{it} = \alpha_i + \lambda_t + \sum_{k \neq -1} \beta_k High_i \mathbf{1}[t - 2020 = k] + \Gamma(X_i \times Post_t) + \varepsilon_{it}. \quad (1)$$

Here α_i are unit fixed effects, λ_t are year fixed effects, and X_i contains the baseline controls (Diemer et al., 2022; Barca, 2009).

The static companion model estimates the average post-period contrast:

$$Y_{it} = \alpha_i + \lambda_t + \delta(High_i \times Post_t) + \Gamma(X_i \times Post_t) + \varepsilon_{it}. \quad (2)$$

The coefficient δ is the compact version of the event path, with standard errors clustered by unit (Card and Krueger, 1994; Autor et al., 2013).

The continuous-exposure row replaces $High_i$ with the pre-period exposure level. That row checks whether the result is a threshold artifact or a gradient across the exposure distribution. The robustness row swaps the outcome to unemployment rate, and the pre-trend row summarizes

whether high-exposure units were already drifting before 2020 (Monte et al., 2018; Redding and Turner, 2015).

4.2 Static Specification

The design supports a capacity-gradient reading. Its strongest evidence comes from agreement among the event path, the static estimate, the continuous-exposure row, and the robustness outcome. Explicit treatment timing, sharper mechanism data, or a credible donor pool would move the design to a stronger causal tier (Donaldson, 2018; Baum-Snow, 2007).

Those stronger tiers are gated by data structure (Rosenbaum, 2002; for National Statistics, 2026). Synthetic difference-in-differences is appropriate when credible donor units and treatment timing are available; partial identification is useful when treatment timing or spillovers make the target parameter interval-valued; double machine learning or causal forest designs are useful when richer covariates can model treatment heterogeneity without replacing the core identification argument (Chodorow-Reich, 2019; Nakamura and Steinsson, 2014).

The omitted event time is $k = -1$, the year immediately before the event (Manski, 2003; Arkhangelsky et al., 2021). That choice gives every plotted coefficient a natural comparison point and prevents the figure from turning into an unanchored sequence of high-versus-low differences. The pre-period coefficients are read as evidence about the comparison, while the post-period coefficients carry the substantive timing result (for National Statistics, 2026; Peng, 2011; Wilkinson et al., 2016).

4.3 Identification

The controls enter as interactions with the post period because they are baseline variables (Imbens and Rubin, 2015; for National Statistics, 2026). A fixed baseline education share or pre-period unemployment rate is absorbed by unit fixed effects, but the interaction lets that baseline characteristic have its own post-period slope. This is the design feature that keeps the exposure coefficient from carrying all baseline advantage by itself (Kitchin, 2014; Borgman, 2015).

The continuous-exposure specification is more than a robustness label. The median split makes the event path readable; the continuous row asks whether additional pre-period exposure is associated with additional movement in employment rate. Agreement between those rows supports a gradient interpretation, while disagreement points toward nonlinearity or threshold effects (Bureau, 2025; Wilkinson et al., 2016).

The robustness outcome is chosen because it is close enough to challenge the mechanism. If unemployment rate moves with employment rate, the interpretation has scope beyond a single dependent variable. If it moves differently, the paper learns that the mechanism is narrower, which is still a useful empirical result (Bartik, 1991; Kline and Moretti, 2014).

Spatial dependence remains a design concern. Unit clustering handles repeated observations within units, while future versions could add spatial covariance estimators, multiway clustering, or boundary-specific designs where neighbouring places are likely to share shocks (Conley, 1999; Cameron et al., 2011; Anselin, 1988; LeSage and Pace, 2009; Gibbons and Overman, 2012).

The identifying comparison is easiest to read as a sequence (Imbens and Rubin, 2015; Callaway and Sant’Anna, 2021). First, the event study asks whether high-exposure units move differently before and after 2020. Second, the static model asks for the average post-period contrast. Third, the continuous-exposure row asks whether the relationship is visible across the exposure distribution. Fourth, the robustness row asks whether the same mechanism appears in unemployment rate (Busso et al., 2013; Greenstone et al., 2010).

Each step has a failure mode (Rosenbaum, 2002). Pre-period drift weakens the timing comparison; a static coefficient with a noisy event path weakens the average contrast; a continuous row that breaks from the median split suggests threshold effects; a robustness row with the opposite sign narrows the mechanism. The paper is designed so these failures are visible rather than hidden in prose (Neumark and Simpson, 2015; Austin et al., 2018).

The method is also sized to the data (Manski, 2003; for National Statistics, 2026). A public annual panel with repeated units supports fixed effects, event-time diagnostics, source-specific controls, and clustered standard errors. It does not need a more ornate estimator unless the additional method answers a data-supported identification problem, such as staggered timing, donor-pool construction, spatial spillovers, or high-dimensional heterogeneity (Moretti, 2011; Glaeser, 2008).

5 Results




Figure 1: NOMIS Annual Population Survey event-study profile

Notes: Points are state-year fixed-effects estimates for high-exposure UK local authority relative to lower-exposure units. Vertical bars are 95% confidence intervals using unit-clustered standard errors. Event time -1 is omitted.

Table 2: Public-panel regression estimates

Specification	Outcome	Coefficient	Estimate	Std. error	95% CI	N
Post-2020 Fixed-effects Contrast with Baseline Controls	Employment Rate	High Exposure by Post-2020	0.463	0.248	[-0.024, 0.949]	3800
Continuous-exposure Post-2020 Contrast with Baseline Controls	Employment Rate	Baseline Exposure by Post-2020	0.045	0.015	[0.015, 0.074]	3800

Notes: Estimates include the fixed effects and clustered standard errors reported in the generated regression CSV.

5.1 Event Study

Figure 1 gives the timing (Imbens and Rubin, 2015; Kitchin, 2014). At event time 5, high-exposure units are lower on employment rate by 1.180 units, with a unit-clustered standard error of 0.608.

The pre-trend placebo coefficient is -0.343 (SE 0.157), which is the first check on whether the post-period contrast is simply an old trajectory continuing (Fajgelbaum and Gaubert, 2020; Gaubert, 2018).

Table 2 gives the compact estimate. The high-exposure post-period coefficient is 0.463 (SE 0.248), so the average post-period contrast is higher for initially high-exposure units. The continuous-exposure row tests the same idea without a median split (Hsieh and Moretti, 2019; Goldsmith-Pinkham et al., 2020).

The robustness row changes the dependent variable to unemployment rate. Its coefficient is -0.134 (SE 0.102). When this row moves with the main outcome, the mechanism has broader support; when it diverges, the paper narrows the interpretation to employment rate instead of overstating the channel (Borusyak et al., 2022; Adao et al., 2019).

The coefficient should be read with the control set in view. It is a within-unit post-period contrast that lets baseline scale, density, education, industry mix, or labour-market attachment have their own post-period slopes (Conley, 1999; Cameron et al., 2011).

5.2 Regression Estimates

The results answer the question at the level the data can support: Do local authorities with higher qualification shares have different employment-rate paths in the APS panel? The evidence is a post-2020 difference in employment rate by pre-2020 NVQ4 qualification share, with timing, magnitude, robustness, and audit files available for inspection.

The event path is read before the static estimate because timing carries economic information. A post-period gap that appears immediately has a different interpretation from one that opens slowly. In this paper, the preferred event-time coefficient is -1.180; the sign, timing, and uncertainty together describe the adjustment path (Duranton and Venables, 2015; Fujita et al., 1999).

The static coefficient then compresses that path into one number. Its value is 0.463, with a standard error of 0.248. The estimate is useful because it is comparable across the paper set: every public panel reports a high-exposure post-period coefficient with the same fixed-effects, controls, and unit-clustering convention (Iammarino et al., 2019; Rodriguez-Pose, 2018).

The continuous-exposure estimate is the scale check. A median split can hide variation inside the high-exposure group, especially when the exposure distribution is skewed. The continuous row asks whether the result survives when the model uses the measured pre-period exposure itself (Diemer et al., 2022; Barca, 2009).

The pre-trend row gives a compact warning light. Its coefficient is -0.343 with a standard error of 0.157. A small and imprecise pre-trend supports the timing comparison; a large pre-trend would shift interpretation toward prior divergence, selection, or the need for a narrower design (Card and Krueger, 1994; Autor et al., 2013).

5.3 Robustness

The robustness estimate is also informative in sign and scale. Its coefficient is -0.134 with a standard error of 0.102. The paper reads that value beside the main estimate rather than using it as a second headline, because the robustness outcome is a test of mechanism scope (Monte et al., 2018; Redding and Turner, 2015).

Precision is interpreted with the data source in mind. Administrative panels can produce tight estimates when the outcome is measured consistently across units; survey panels can be noisier; county migration rates can be sensitive to population denominators. The standard errors therefore

guide inference, while the coefficient path guides economic reading (Donaldson, 2018; Baum-Snow, 2007).

The result also identifies the next empirical bottleneck. When the coefficient is clear but the mechanism remains broad, the next evidence should measure the missing channel directly. For this paper, that means: A stronger design would combine APS with vacancies, sector shares, commuting flows, or administrative employment data (Chodorow-Reich, 2019; Nakamura and Steinsson, 2014)..

The summary table matters for interpreting the estimates (Kline and Moretti, 2014; Gelman et al., 2020). It reports the number of units, years, retained rows, exposure threshold, and controls. Those quantities tell the reader whether the coefficient is coming from a broad panel or a narrow slice, and whether the control set has enough economic content to support the stated comparison (Fajgelbaum and Gaubert, 2020; Gaubert, 2018).

The event-study table in the appendix is the numerical companion to the figure. It lets the reader inspect each event time, confidence interval, sample size, fixed-effect structure, and standard-error convention. That is important because a plotted path should be beautiful enough to read quickly and precise enough to audit slowly (Hsieh and Moretti, 2019; Goldsmith-Pinkham et al., 2020).

The regression table is the compact comparison across designs. The high-exposure row gives the headline contrast, while the continuous-exposure row keeps scale in view. Reading both rows prevents the paper from leaning too heavily on the median cutoff, especially when the exposure distribution has a long upper tail (Borusyak et al., 2022; Adao et al., 2019).

The robustness table is the place where the paper earns restraint. It does not erase the main result when the neighbouring outcome is weaker, and it does not inflate the result when the neighbouring outcome agrees. It tells the reader how far the mechanism travels inside the available data (Conley, 1999; Cameron et al., 2011).

Taken as a package, the results have a clear interpretation path. Start with the event figure for timing, move to the regression table for average magnitude, check the continuous row for scale, and read the robustness row for mechanism scope (Angrist and Pischke, 2009; Wooldridge, 2010).

The data trail is part of the result, not an appendix afterthought. The retained CSVs and manifest let the reader move from the coefficient table back to the underlying unit-year panel, check the exposure threshold, and verify which controls entered the post-period interactions (Kitchin, 2014; Wilkinson et al., 2016; Peng, 2011).

The result should also be read against the size of the panel (Moretti, 2011; Rosenbaum, 2002). With 380 units and 3,800 unit-year observations, the model has enough repeated local variation to estimate a fixed-effects contrast, but the interpretation still depends on source definitions and the control set. That balance is what makes the estimate useful for screening the next design layer (Imbens and Rubin, 2015; Rosenbaum, 2002).

The final empirical reading is therefore neither a bare sign nor a significance label (Rosenbaum, 2002; Kitchin, 2014). It is the combination of a measured exposure, a local outcome, timing evidence, a compact post-period coefficient, a neighbouring robustness outcome, and a data audit. A future paper can replace pieces of that package with stronger data, but it should preserve the same interpretive discipline (Manski, 2003; Rubin, 1974).

6 Discussion

6.1 Interpretation

Qualification structure marks a labour-supply capacity margin. The estimates translate that idea into a visible local adjustment pattern: high-exposure units sit higher on the post-period employment rate contrast, while the event path shows when the gap opens or closes.

For this APS paper, the coefficient is about human-capital composition and employment-rate adjustment. A positive estimate means high-qualification places sustain stronger employment-rate paths; a negative estimate means qualifications are not translating into stronger local employment-rate adjustment in this panel.

The important criticism is residual composition. Sectoral exposure, housing constraints, commuting links, local institutions, and programme timing can still sit behind the exposure. The current controls reduce obvious baseline pressure; they also show exactly which variables a stronger next design should add ([Angrist and Pischke, 2009](#); [Wooldridge, 2010](#)).

A stronger design would combine APS with vacancies, sector shares, commuting flows, or administrative employment data. That evidence would move the design from a capacity-gradient estimate toward a mechanism-specific comparison.

6.2 Threats

The result is therefore useful for deciding what evidence matters next. It identifies whether pre-2020 NVQ4 qualification share deserves sharper treatment data, richer covariates, or a mechanism-specific design.

The interpretation is strongest when the main estimate, the event path, and the robustness row point in a compatible direction. That pattern says the measured exposure is organizing the local outcome, even before a richer policy or mechanism file is joined to the panel ([Callaway and Sant’Anna, 2021](#); [Sun and Abraham, 2021](#)).

The criticism is useful because it is specific. The current paper needs better information on employment adjustment through local human-capital composition, neighbouring-place spillovers, and explicit timing for policies or shocks. Those needs are narrower than a generic call for better data, and they point to a concrete follow-up design ([Goodman-Bacon, 2021](#); [Arkhangelsky et al., 2021](#)).

Policy reading should therefore follow the measured channel ([Abadie et al., 2010](#); [Athey and Imbens, 2016](#)). A result about employment rate should guide measurement of employment rate and its nearby mechanisms; welfare, multiplier, or programme-effect questions need their own treatment files and mechanism data ([Chernozhukov et al., 2018](#); [Wager and Athey, 2018](#)).

Weaker local-adjustment designs improve when they change the empirical object, add source-appropriate controls, lengthen the panel where the data allow it, and make the next missing file explicit ([Athey and Wager, 2019](#); [Mullainathan and Spiess, 2017](#)).

The remaining uncertainty is productive ([Gentzkow et al., 2019](#); [Grimmer and Stewart, 2013](#)). It tells the next study what to look for: a treatment date, a mechanism measure, a spatial link, a richer covariate set, or a more precise outcome. The present paper therefore functions as both an estimate and an evidence map for the next design layer ([Gelman and Hill, 2007](#); [Gelman et al., 2020](#)).

The discussion also clarifies what would raise the standard ([Kitchin, 2014](#); [Borgman, 2015](#)). Stronger papers come from a better match among question, source, outcome, exposure, controls, method, and interpretation ([Anselin, 1988](#); [LeSage and Pace, 2009](#)).

The comparison with benchmark papers is demanding but useful. Those papers earn broad claims through stronger shocks, richer administrative data, sharper comparison groups, or explicit mechanisms. This paper earns a narrower claim by making the public-data comparison transparent and by showing which additional file would raise the design tier ([Openshaw, 1984](#); [Fotheringham et al., 2002](#)).

A second reading concerns scale ([Bartik, 1991](#); [Manski, 2003](#)). The coefficient is not just a sign on a table; it describes how much the outcome differs when the pre-period local margin is high, after

the model absorbs permanent unit differences, common year shocks, and baseline controls. That scale helps distinguish economically visible adjustment from a statistically tidy but substantively small contrast (Pearl, 2009; Keele, 2015).

The controls also sharpen the economic reading (Borgman, 2015). Log population addresses market size, density addresses spatial pressure where it is not itself the exposure, education or qualifications address human-capital composition, and labour-market controls address pre-existing slack. Reading the estimate alongside those controls makes the comparison closer to the local-adjustment arguments in the literature (King et al., 1995; Shadish et al., 2002).

That is why the robustness row matters. It asks whether the local mechanism appears only in the headline outcome or also in a neighbouring margin with related economic content. Agreement widens the interpretation; disagreement tells the reader where the mechanism stops (Peng, 2011; Wilkinson et al., 2016).

The central improvement is that the estimate now has friction. The result can be challenged through the pre-trend row, the continuous-exposure row, the robustness outcome, the control interactions, and the audit file. Those objects create places where a reviewer can disagree productively (for National Statistics, 2026; Peng, 2011; Wilkinson et al., 2016).

The same principle should govern related questions (Borgman, 2015). A comparison should be narrowed if its panel is too short, its controls duplicate the exposure, its prose leans on broad post-event language, or its conclusion spends more time disclaiming than interpreting. Stronger papers make a measured claim, state the economic mechanism, and expose the next empirical bottleneck (Kitchin, 2014; Borgman, 2015).

The broader substantive lesson is simple: the claim improves only when the data, method, and prose all become more specific. For this paper, specificity comes from the 2016–2025 panel, the pre-2020 NVQ4 qualification share exposure, the employment rate outcome, the unemployment rate robustness check, and the stated mechanism upgrade (Bureau, 2025; Wilkinson et al., 2016).

The practical criterion is direct: a reader should be able to name the empirical object in one sentence, inspect the derived panel, follow the coefficient path, and state the next data join (Chodorow-Reich, 2019; Imbens and Rubin, 2015). If any of those steps fails, the interpretation should be narrowed at the point where the manuscript reads the result (Neumark and Simpson, 2015; Austin et al., 2018).

6.3 Conclusion

The paper contributes a separate empirical object: NOMIS Annual Population Survey (Imbens and Rubin, 2015; Rosenbaum, 2002), employment rate, pre-2020 NVQ4 qualification share, and controls for pre-2020 employment rate, pre-2020 economic activity rate, pre-2020 unemployment rate. It answers the stated question with a fixed-effects event path, a compact regression estimate, a continuous-exposure check, and a robustness outcome.

The preferred reading follows the evidence. The sign and timing of the estimates describe how the measured local capacity margin is related to the observed outcome; the control set and pre-trend check determine how much confidence to place in that reading.

Interpreted through observational-design standards (Manski, 2003; Rubin, 1974), aPS local-authority estimates are survey based, so small-area noise and changing samples matter for interpretation. The next empirical step is also clear: A stronger design would combine APS with vacancies, sector shares, commuting flows, or administrative employment data.

The evidence is most persuasive when the panel, estimates, and prose all tell the same story (Bartik, 1991; Kline and Moretti, 2014): the source is public, the comparison is explicit, the controls are motivated by the literature, and the remaining evidence need is precise.

For this question, that story is now specific ([Busso et al., 2013](#); [Greenstone et al., 2010](#)): NOMIS Annual Population Survey measures employment rate, pre-2020 NVQ4 qualification share defines the pre-period capacity margin, and the controls for pre-2020 employment rate, pre-2020 economic activity rate, pre-2020 unemployment rate keep the comparison anchored in standard local-labour-market concerns. The next empirical step is no longer vague; it is the mechanism join named above.

The paper therefore sits in the applied-economics tradition of matching local mechanisms to credible comparisons, while using fixed effects and event-study diagnostics as the minimum evidentiary baseline ([Moretti, 2011](#); [Greenstone et al., 2010](#); [Angrist and Pischke, 2009](#); [Wooldridge, 2010](#)).

A Data Construction

The data construction appendix reports the derived public panel and its provenance. Raw public-data payloads are hashed and discarded after the derived files are written (Peng, 2011; Wilkinson et al., 2016).

B Supplementary Source Evidence

The supplementary source evidence records the retrieval surface used for framing and reference checks. It keeps the literature sample, access rates, and search diagnostics visible without substituting them for the subject data (Priem et al., 2022; OpenAlex, 2026).

Table 3: Open-access literature sample

Quantity	Value
Retained open-access records	53
Records with direct document link	53
Document-access share	1.000
Publication-year cells	5

Table 4: Evidence-coverage regression estimates

Model	Outcome	Coef.	SE	95% CI	N
Annual open-access record trend	records per year	-2.100	0.709	[-3.491, -0.709]	5
PDF availability trend	direct PDF indicator	0.000	0.000	[0.000, 0.000]	53
Topic relevance and attention	log cited-by count	-15.603	6.834	[-28.999, -2.208]	53

Table 5: Robustness checks for the evidence-coverage trend

Sample	N	Trend	SE	95% CI
all records	53	-2.100	0.709	[-3.491, -0.709]
high relevance	2	0.000	0.000	[0.000, 0.000]
recent records	14	0.000	0.000	[0.000, 0.000]
PDF available	53	-2.100	0.709	[-3.491, -0.709]

C Supporting Regressions and Robustness

The supporting estimates expand the main result into the event-study table, robustness outcome, and pre-trend placebo. Baseline controls are interacted with the post period so the comparison is not carried by the exposure variable alone (Angrist and Pischke, 2009; Wooldridge, 2010).

Table 6: NOMIS Annual Population Survey event-study estimates

Model	Outcome	Event time	Coefficient	Std. error	95% CI	N
UK Local Authority Event Study	Employment Rate	-4	0.827	0.472	[-0.098, 1.753]	3800
UK Local Authority Event Study	Employment Rate	-3	0.959	0.467	[0.044, 1.874]	3800
UK Local Authority Event Study	Employment Rate	-2	0.008	0.412	[-0.800, 0.817]	3800
UK Local Authority Event Study	Employment Rate	0	-0.255	0.445	[-1.128, 0.618]	3800
UK Local Authority Event Study	Employment Rate	1	1.641	0.676	[0.316, 2.965]	3800
UK Local Authority Event Study	Employment Rate	2	1.400	0.502	[0.416, 2.384]	3800
UK Local Authority Event Study	Employment Rate	3	1.335	0.566	[0.226, 2.445]	3800
UK Local Authority Event Study	Employment Rate	4	0.656	0.670	[-0.656, 1.969]	3800
UK Local Authority Event Study	Employment Rate	5	-1.180	0.608	[-2.372, 0.012]	3800

Notes: Event time -1 is the omitted category. Estimates include unit and year fixed effects with unit-clustered standard errors.

Table 7: Public-panel robustness checks

Robustness check	Outcome	Coefficient	Estimate	Std. error	95% CI	N
Alternative outcome: unemployment rate	Unemployment rate	High Exposure by Post-2020	-0.134	0.102	[-0.333, 0.065]	3800
Pre-trend placebo	Employment Rate	High Exposure by Pre-period Trend	-0.343	0.157	[-0.651, -0.036]	1520

Notes: Estimates include the fixed effects and clustered standard errors reported in the generated regression CSV.

References

- Alberto Abadie, Alexis Diamond, and Jens Hainmueller. Synthetic control methods for comparative case studies. *Journal of the American Statistical Association*, 2010.
- Rodrigo Adao, Michal Kolesar, and Eduardo Morales. Shift-share designs: Theory and inference. *Quarterly Journal of Economics*, 2019.
- Joshua D. Angrist and Joern-Steffen Pischke. Mostly harmless econometrics. Princeton University Press, 2009.
- Luc Anselin. Spatial econometrics: Methods and models. Kluwer Academic Publishers, 1988.
- Dmitry Arkhangelsky, Susan Athey, David A. Hirshberg, Guido W. Imbens, and Stefan Wager. Synthetic difference-in-differences. *American Economic Review*, 2021.
- Susan Athey and Guido W. Imbens. Recursive partitioning for heterogeneous causal effects. *Proceedings of the National Academy of Sciences*, 2016.
- Susan Athey and Stefan Wager. Estimating treatment effects with causal forests. *Observational Studies*, 2019.
- Benjamin Austin, Edward Glaeser, and Lawrence H. Summers. Jobs for the heartland: Place-based policies in 21st-century america. *Brookings Papers on Economic Activity*, 2018.
- David H. Autor, David Dorn, and Gordon H. Hanson. The china syndrome. *American Economic Review*, 2013.
- Fabrizio Barca. An agenda for a reformed cohesion policy. European Commission report, 2009.
- Timothy J. Bartik. Who benefits from state and local economic development policies? W.E. Upjohn Institute, 1991.
- Nathaniel Baum-Snow. Did highways cause suburbanization? *Quarterly Journal of Economics*, 2007.
- Christine L. Borgman. Big data, little data, no data. MIT Press, 2015.
- Kirill Borusyak, Peter Hull, and Xavier Jaravel. Quasi-experimental shift-share research designs. *Review of Economic Studies*, 2022.
- United States Census Bureau. Business dynamics statistics bdsgeo geography file. Census Bureau public data file, 2025.
- Matias Busso, Jesse Gregory, and Patrick Kline. Assessing the incidence and efficiency of a prominent place based policy. *American Economic Review*, 2013.
- Brantly Callaway and Pedro H. C. Sant’Anna. Difference-in-differences with multiple time periods. *Journal of Econometrics*, 2021.
- A. Colin Cameron, Jonah B. Gelbach, and Douglas L. Miller. Robust inference with multiway clustering. *Journal of Business and Economic Statistics*, 2011.
- David Card and Alan B. Krueger. Minimum wages and employment: A case study. *American Economic Review*, 1994.

- Victor Chernozhukov, Denis Chetverikov, Mert Demirer, Esther Duflo, Christian Hansen, Whitney Newey, and James Robins. Double/debiased machine learning for treatment and structural parameters. *The Econometrics Journal*, 2018.
- Gabriel Chodorow-Reich. Geographic cross-sectional fiscal spending multipliers: What have we learned? *American Economic Journal: Economic Policy*, 2019.
- Timothy G. Conley. Gmm estimation with cross sectional dependence. *Journal of Econometrics*, 1999.
- Andreas Diemer, Simona Iammarino, Andres Rodriguez-Pose, and Michael Storper. The regional development trap in europe. *Economic Geography*, 2022.
- Dave Donaldson. Railroads of the raj: Estimating the impact of transportation infrastructure. *American Economic Review*, 2018.
- Gilles Duranton and Anthony J. Venables. Place-based policies for development. *Handbook of Regional and Urban Economics*, 2015.
- Pablo D. Fajgelbaum and Cecile Gaubert. Optimal spatial policies, geography, and sorting. *Quarterly Journal of Economics*, 2020.
- Office for National Statistics. Nomis annual population survey. NOMIS public data API, 2026.
- A. Stewart Fotheringham, Chris Brunsdon, and Martin Charlton. Geographically weighted regression. Wiley, 2002.
- Masahisa Fujita, Paul Krugman, and Anthony J. Venables. The spatial economy. MIT Press, 1999.
- Cecile Gaubert. Firm sorting and agglomeration. *American Economic Review*, 2018.
- Andrew Gelman and Jennifer Hill. Data analysis using regression and multilevel/hierarchical models. Cambridge University Press, 2007.
- Andrew Gelman, Aki Vehtari, Daniel Simpson, Charles C. Margossian, Bob Carpenter, Yuling Yao, Lauren Kennedy, Jonah Gabry, Paul-Christian Buerkner, and Martin Modrak. Bayesian workflow. arXiv preprint, 2020.
- Matthew Gentzkow, Bryan Kelly, and Matt Taddy. Text as data. *Journal of Economic Literature*, 2019.
- Stephen Gibbons and Henry G. Overman. Mostly pointless spatial econometrics? *Journal of Regional Science*, 2012.
- Edward L. Glaeser. Cities, agglomeration, and spatial equilibrium. Oxford University Press, 2008.
- Paul Goldsmith-Pinkham, Isaac Sorkin, and Henry Swift. Bartik instruments: What, when, why, and how. *American Economic Review*, 2020.
- Andrew Goodman-Bacon. Difference-in-differences with variation in treatment timing. *Journal of Econometrics*, 2021.
- Michael Greenstone, Richard Hornbeck, and Enrico Moretti. Identifying agglomeration spillovers: Evidence from winners and losers of large plant openings. *Journal of Political Economy*, 2010.

- Justin Grimmer and Brandon M. Stewart. Text as data: The promise and pitfalls. *Political Analysis*, 2013.
- Chang-Tai Hsieh and Enrico Moretti. Housing constraints and spatial misallocation. *American Economic Journal: Macroeconomics*, 2019.
- Simona Iammarino, Andres Rodriguez-Pose, and Michael Storper. Regional inequality in europe: Evidence, theory and policy implications. *Journal of Economic Geography*, 2019.
- Guido W. Imbens and Donald B. Rubin. Causal inference for statistics, social, and biomedical sciences. Cambridge University Press, 2015.
- Luke Keele. The statistics of causal inference: A view from political methodology. *Political Analysis*, 2015.
- Gary King, Robert O. Keohane, and Sidney Verba. Designing social inquiry. Princeton University Press, 1995.
- Rob Kitchin. The data revolution. SAGE, 2014.
- Patrick Kline and Enrico Moretti. People, places, and public policy: Some simple welfare economics of local economic development programs. *Annual Review of Economics*, 2014.
- James LeSage and R. Kelley Pace. Introduction to spatial econometrics. CRC Press, 2009.
- Charles F. Manski. Partial identification of probability distributions. Springer, 2003.
- Ferdinando Monte, Stephen J. Redding, and Esteban Rossi-Hansberg. Commuting, migration, and local employment elasticities. *American Economic Review*, 2018.
- Enrico Moretti. Local labor markets. *Handbook of Labor Economics*, 2011.
- Sendhil Mullainathan and Jann Spiess. Machine learning: An applied econometric approach. *Journal of Economic Perspectives*, 2017.
- Emi Nakamura and Jon Steinsson. Fiscal stimulus in a monetary union: Evidence from us regions. *American Economic Review*, 2014.
- David Neumark and Helen Simpson. Place-based policies. *Handbook of Regional and Urban Economics*, 2015.
- OpenAlex. Openalex works api documentation. OpenAlex documentation, 2026.
- Stan Openshaw. The modifiable areal unit problem. Geo Books, 1984.
- Judea Pearl. Causality. Cambridge University Press, 2009.
- Roger D. Peng. Reproducible research in computational science. *Science*, 2011.
- Jason Priem, Heather Piwowar, and Richard Orr. Openalex: A fully-open index of scholarly works, authors, venues, institutions, and concepts. arXiv preprint, 2022.
- Stephen J. Redding and Matthew A. Turner. Transportation costs and the spatial organization of economic activity. *Handbook of Regional and Urban Economics*, 2015.

- Andres Rodriguez-Pose. The revenge of the places that don't matter. Cambridge Journal of Regions, Economy and Society, 2018.
- Paul R. Rosenbaum. Observational studies. Springer, 2002.
- Donald B. Rubin. Estimating causal effects of treatments in randomized and nonrandomized studies. Journal of Educational Psychology, 1974.
- William R. Shadish, Thomas D. Cook, and Donald T. Campbell. Experimental and quasi-experimental designs. Houghton Mifflin, 2002.
- Liyang Sun and Sarah Abraham. Estimating dynamic treatment effects in event studies with heterogeneous treatment effects. Journal of Econometrics, 2021.
- Stefan Wager and Susan Athey. Estimation and inference of heterogeneous treatment effects using random forests. Journal of the American Statistical Association, 2018.
- Mark D. Wilkinson, Michel Dumontier, IJsbrand Jan Aalbersberg, et al. The fair guiding principles. Scientific Data, 2016.
- Jeffrey M. Wooldridge. Econometric analysis of cross section and panel data. MIT Press, 2010.